

Convolutional neural network layers

In this notebook, we will build the convolutional neural network layers. This will be followed by a spatial batchnorm, and then in the final notebook of this assignment, we will train a CNN to further improve the validation accuracy on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

In [1]:

```
## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

Implementing CNN layers

Just as we implemented modular layers for fully connected networks, batch normalization, and dropout, we'll want to implement modular layers for convolutional neural networks. These layers are in `nndl/conv_layers.py`.

Convolutional forward pass

Begin by implementing a naive version of the forward pass of the CNN that uses for loops. This function is `conv_forward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a triple for loop.

After you implement `conv_forward_naive`, test your implementation by running the cell below.

In [8]:

```
x_shape = (2, 3, 4, 4)
w_shape = (3, 3, 4, 4)
x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
w = np.linspace(-0.2, 0.3, num=np.prod(w_shape)).reshape(w_shape)
b = np.linspace(-0.1, 0.2, num=3)

conv_param = {'stride': 2, 'pad': 1}
out, _ = conv_forward_naive(x, w, b, conv_param)
correct_out = np.array([[[[-0.08759809, -0.10987781],
                           [-0.18387192, -0.2109216 ]],
                          [[ 0.21027089,  0.21661097],
                           [ 0.22847626,  0.23004637]],
                          [[ 0.50813986,  0.54309974],
                           [ 0.64082444,  0.67101435]]],
                         [[[-0.98053589, -1.03143541],
                           [-1.19128892, -1.24695841]],
                          [[ 0.69108355,  0.66880383],
                           [ 0.59480972,  0.56776003]],
                          [[ 2.36270298,  2.36904306],
                           [ 2.38090835,  2.38247847]]]])

# Compare your output to ours; difference should be around 1e-8
print('Testing conv_forward_naive')
print('difference: ', rel_error(out, correct_out))
```

```
Testing conv_forward_naive
difference: 2.2121476417505994e-08
```

Convolutional backward pass

Now, implement a naive version of the backward pass of the CNN. The function is `conv_backward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a quadruple for loop.

After you implement `conv_backward_naive`, test your implementation by running the cell below.

In [11]:

```
x = np.random.randn(4, 3, 5, 5)
w = np.random.randn(2, 3, 3, 3)
b = np.random.randn(2,)
dout = np.random.randn(4, 2, 5, 5)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_forward_naive(x,w,b,conv_param)

dx_num = eval_numerical_gradient_array(lambda x: conv_forward_naive(x, w, b, conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_forward_naive(x, w, b, conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_forward_naive(x, w, b, conv_param)[0], b, dout)

out, cache = conv_forward_naive(x, w, b, conv_param)
dx, dw, db = conv_backward_naive(dout, cache)

# Your errors should be around 1e-9'
print('Testing conv_backward_naive function')
print('dx error: ', rel_error(dx, dx_num))
print('dw error: ', rel_error(dw, dw_num))
print('db error: ', rel_error(db, db_num))
```

Testing conv_backward_naive function

dx error: 2.6751749859316384e-09

dw error: 5.287927760897462e-10

db error: 5.297964130100186e-10

Max pool forward pass

In this section, we will implement the forward pass of the max pool. The function is `max_pool_forward_naive` in `nnd1/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_forward_naive`, test your implementation by running the cell below.

In [13]:

```
x_shape = (2, 3, 4, 4)
x = np.linspace(-0.3, 0.4, num=np.prod(x_shape)).reshape(x_shape)
pool_param = {'pool_width': 2, 'pool_height': 2, 'stride': 2}

out, _ = max_pool_forward_naive(x, pool_param)

correct_out = np.array([[[[-0.26315789, -0.24842105],
                           [-0.20421053, -0.18947368]],
                          [[-0.14526316, -0.13052632],
                           [-0.08631579, -0.07157895]],
                          [[-0.02736842, -0.01263158],
                           [ 0.03157895,  0.04631579]]],
                        [[[ 0.09052632,  0.10526316],
                           [ 0.14947368,  0.16421053]],
                         [[ 0.20842105,  0.22315789],
                           [ 0.26736842,  0.28210526]],
                         [[ 0.32631579,  0.34105263],
                           [ 0.38526316,  0.4          ]]]]])

# Compare your output with ours. Difference should be around 1e-8.
print('Testing max_pool_forward_naive function:')
print('difference: ', rel_error(out, correct_out))
```

Testing max_pool_forward_naive function:
 difference: 4.1666665157267834e-08

Max pool backward pass

In this section, you will implement the backward pass of the max pool. The function is `max_pool_backward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_backward_naive`, test your implementation by running the cell below.

In [14]:

```
x = np.random.randn(3, 2, 8, 8)
dout = np.random.randn(3, 2, 4, 4)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

dx_num = eval_numerical_gradient_array(lambda x: max_pool_forward_naive(x, pool_param)[0], x, dout)

out, cache = max_pool_forward_naive(x, pool_param)
dx = max_pool_backward_naive(dout, cache)

# Your error should be around 1e-12
print('Testing max_pool_backward_naive function:')
print('dx error: ', rel_error(dx, dx_num))
```

Testing max_pool_backward_naive function:
 dx error: 3.275641145949778e-12

Fast implementation of the CNN layers

Implementing fast versions of the CNN layers can be difficult. We will provide you with the fast layers implemented by cs231n. They are provided in `cs231n/fast_layers.py`.

The fast convolution implementation depends on a Cython extension; to compile it you need to run the following from the `cs231n` directory:

```
python setup.py build_ext --inplace
```

NOTE: The fast implementation for pooling will only perform optimally if the pooling regions are non-overlapping and tile the input. If these conditions are not met then the fast pooling implementation will not be much faster than the naive implementation.

You can compare the performance of the naive and fast versions of these layers by running the cell below.

You should see pretty drastic speedups in the implementation of these layers. On our machine, the forward pass speeds up by 17x and the backward pass speeds up by 840x. Of course, these numbers will vary from machine to machine, as well as on your precise implementation of the naive layers.

In [17]:

```
from cs231n.fast_layers import conv_forward_fast, conv_backward_fast
from time import time

x = np.random.randn(100, 3, 31, 31)
w = np.random.randn(25, 3, 3, 3)
b = np.random.randn(25,)
dout = np.random.randn(100, 25, 16, 16)
conv_param = {'stride': 2, 'pad': 1}

t0 = time()
out_naive, cache_naive = conv_forward_naive(x, w, b, conv_param)
t1 = time()
out_fast, cache_fast = conv_forward_fast(x, w, b, conv_param)
t2 = time()

print('Testing conv_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('Fast: %fs' % (t2 - t1))
print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('Difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive, dw_naive, db_naive = conv_backward_naive(dout, cache_naive)
t1 = time()
dx_fast, dw_fast, db_fast = conv_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting conv_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('Fast: %fs' % (t2 - t1))
print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))
print('dw difference: ', rel_error(dw_naive, dw_fast))
print('db difference: ', rel_error(db_naive, db_fast))
```

```
Testing conv_forward_fast:
Naive: 8.299330s
Fast: 0.032375s
Speedup: 256.352952x
Difference: 7.834443839149003e-12
```

```
Testing conv_backward_fast:
Naive: 17.086036s
Fast: 0.015629s
Speedup: 1093.239398x
dx difference: 2.095174024872286e-10
dw difference: 4.781055692493398e-13
db difference: 6.658772224755541e-15
```

In [19]:

```
from cs231n.fast_layers import max_pool_forward_fast, max_pool_backward_fast

x = np.random.randn(100, 3, 32, 32)
dout = np.random.randn(100, 3, 16, 16)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

t0 = time()
out_naive, cache_naive = max_pool_forward_naive(x, pool_param)
t1 = time()
out_fast, cache_fast = max_pool_forward_fast(x, pool_param)
t2 = time()

print('Testing pool_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('fast: %fs' % (t2 - t1))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive = max_pool_backward_naive(dout, cache_naive)
t1 = time()
dx_fast = max_pool_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting pool_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1) + 1e-7))
print('dx difference: ', rel_error(dx_naive, dx_fast))
```

```
Testing pool_forward_fast:
Naive: 0.611588s
fast: 0.005098s
speedup: 119.969367x
difference: 0.0
```

```
Testing pool_backward_fast:
Naive: 2.217211s
speedup: 67.174146x
dx difference: 0.0
```

Implementation of cascaded layers

We've provided the following functions in `nndl/conv_layer_utils.py`:

- `conv_relu_forward`
- `conv_relu_backward`
- `conv_relu_pool_forward`
- `conv_relu_pool_backward`

These use the fast implementations of the conv net layers. You can test them below:

In [20]:

```
from nndl.conv_layer_utils import conv_relu_pool_forward, conv_relu_pool_backward

x = np.random.randn(2, 3, 16, 16)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

out, cache = conv_relu_pool_forward(x, w, b, conv_param, pool_param)
dx, dw, db = conv_relu_pool_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_pool_forward(x, w, b, conv_param, pool_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_pool_forward(x, w, b, conv_param, pool_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_pool_forward(x, w, b, conv_param, pool_param)[0], b, dout)

print('Testing conv_relu_pool')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

```
Testing conv_relu_pool
dx error: 1.4687682629736857e-08
dw error: 7.081563484375007e-09
db error: 5.722554013219082e-11
```

In [21]:

```
from nndl.conv_layer_utils import conv_relu_forward, conv_relu_backward

x = np.random.randn(2, 3, 8, 8)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_relu_forward(x, w, b, conv_param)
dx, dw, db = conv_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_forward(x, w, b, conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_forward(x, w, b, conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_forward(x, w, b, conv_param)[0], b, dout)

print('Testing conv_relu:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

```
Testing conv_relu:
dx error: 1.935760425875706e-09
dw error: 2.2368672004727702e-10
db error: 6.883358128649818e-11
```


What next?

We saw how helpful batch normalization was for training FC nets. In the next notebook, we'll implement a batch normalization for convolutional neural networks, and then finish off by implementing a CNN to improve our validation accuracy on CIFAR-10.

Spatial batch normalization

In fully connected networks, we performed batch normalization on the activations. To do something equivalent on CNNs, we modify batch normalization slightly.

Normally batch-normalization accepts inputs of shape (N, D) and produces outputs of shape (N, D) , where we normalize across the minibatch dimension N . For data coming from convolutional layers, batch normalization accepts inputs of shape (N, C, H, W) and produces outputs of shape (N, C, H, W) where the N dimension gives the minibatch size and the (H, W) dimensions give the spatial size of the feature map.

How do we calculate the spatial averages? First, notice that for the C feature maps we have (i.e., the layer has C filters) that each of these ought to have its own batch norm statistics, since each feature map may be picking out very different features in the images. However, within a feature map, we may assume that across all inputs and across all locations in the feature map, there ought to be relatively similar first and second order statistics. Hence, one way to think of spatial batch-normalization is to reshape the (N, C, H, W) array as an $(N * H * W, C)$ array and perform batch normalization on this array.

Since spatial batch norm and batch normalization are similar, it'd be good to at this point also copy and paste our prior implemented layers from HW #4. Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4:

- `layers.py` for your FC network layers, as well as `batchnorm` and `dropout`.
- `layer_utils.py` for your combined FC network layers.
- `optim.py` for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

If you use your prior implementations of the `batchnorm`, then your `spatial batchnorm` implementation may be very short. Our implementations of the forward and backward pass are each 6 lines of code.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

In [2]:

```
## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

Spatial batch normalization forward pass

Implement the forward pass, `spatial_batchnorm_forward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

In [3]:

```
# Check the training-time forward pass by checking means and variances
# of features both before and after spatial batch normalization
```

```
N, C, H, W = 2, 3, 4, 5
x = 4 * np.random.randn(N, C, H, W) + 10

print('Before spatial batch normalization:')
print('  Shape: ', x.shape)
print('  Means: ', x.mean(axis=(0, 2, 3)))
print('  Stds: ', x.std(axis=(0, 2, 3)))

# Means should be close to zero and stds close to one
gamma, beta = np.ones(C), np.zeros(C)
bn_param = {'mode': 'train'}
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization:')
print('  Shape: ', out.shape)
print('  Means: ', out.mean(axis=(0, 2, 3)))
print('  Stds: ', out.std(axis=(0, 2, 3)))

# Means should be close to beta and stds close to gamma
gamma, beta = np.asarray([3, 4, 5]), np.asarray([6, 7, 8])
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization (nontrivial gamma, beta):')
print('  Shape: ', out.shape)
print('  Means: ', out.mean(axis=(0, 2, 3)))
print('  Stds: ', out.std(axis=(0, 2, 3)))
```

```
Before spatial batch normalization:
  Shape: (2, 3, 4, 5)
  Means: [ 9.57416605  9.27706117 10.1095548 ]
  Stds: [3.58436747 3.85186965 4.59523291]
After spatial batch normalization:
  Shape: (2, 3, 4, 5)
  Means: [-3.69149156e-16  2.66453526e-16  2.22044605e-17]
  Stds: [0.99999961 0.99999966 0.99999976]
After spatial batch normalization (nontrivial gamma, beta):
  Shape: (2, 3, 4, 5)
  Means: [6. 7. 8.]
  Stds: [2.99999883 3.99999865 4.99999882]
```

Spatial batch normalization backward pass

Implement the backward pass, `spatial_batchnorm_backward` in `nnd1/conv_layers.py`. Test your implementation by running the cell below.

In [5]:

```
N, C, H, W = 2, 3, 4, 5
x = 5 * np.random.randn(N, C, H, W) + 12
gamma = np.random.randn(C)
beta = np.random.randn(C)
dout = np.random.randn(N, C, H, W)

bn_param = {'mode': 'train'}
fx = lambda x: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fg = lambda a: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fb = lambda b: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]

dx_num = eval_numerical_gradient_array(fx, x, dout)
da_num = eval_numerical_gradient_array(fg, gamma, dout)
db_num = eval_numerical_gradient_array(fb, beta, dout)

_, cache = spatial_batchnorm_forward(x, gamma, beta, bn_param)
dx, dgamma, dbeta = spatial_batchnorm_backward(dout, cache)
print('dx error: ', rel_error(dx_num, dx))
print('dgamma error: ', rel_error(da_num, dgamma))
print('dbeta error: ', rel_error(db_num, dbeta))
```

```
dx error:  1.1144232888574363e-08
dgamma error:  5.228123538106616e-12
dbeta error:  3.275637824095975e-12
```

Convolutional neural networks

In this notebook, we'll put together our convolutional layers to implement a 3-layer CNN. Then, we'll ask you to implement a CNN that can achieve $> 65\%$ validation error on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

If you have not completed the Spatial BatchNorm Notebook, please see the following description from that notebook:

Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their layer implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4:

- `layers.py` for your FC network layers, as well as batchnorm and dropout.
- `layer_utils.py` for your combined FC network layers.
- `optim.py` for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

In [1]:

```
# As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
from nndl.cnn import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient_array, eval_numerical_gradient
from nndl.layers import *
from nndl.conv_layers import *
from cs231n.fast_layers import *
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

In [2]:

```
# Load the (preprocessed) CIFAR10 data.

data = get_CIFAR10_data()
for k in data.keys():
    print('{}: {}'.format(k, data[k].shape))

X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

Three layer CNN

In this notebook, you will implement a three layer CNN. The `ThreeLayerConvNet` class is in `nnd1/cnn.py`. You'll need to modify that code for this section, including the initialization, as well as the calculation of the loss and gradients. You should be able to use the building blocks you have either earlier coded or that we have provided. Be sure to use the fast layers.

The architecture of this CNN will be:

conv - relu - 2x2 max pool - affine - relu - affine - softmax

We won't use batchnorm yet. You've also done enough of these to know how to debug; use the cells below.

Note: As we are implementing several layers CNN networks. The gradient error can be expected for the `eval_numerical_gradient()` function. If your `W1` max relative error and `W2` max relative error are around or below 0.01, they should be acceptable. Other errors should be less than $1e-5$.

In [4]:

```
num_inputs = 2
input_dim = (3, 16, 16)
reg = 0.0
num_classes = 10
X = np.random.randn(num_inputs, *input_dim)
y = np.random.randint(num_classes, size=num_inputs)

model = ThreeLayerConvNet(num_filters=3, filter_size=3,
                           input_dim=input_dim, hidden_dim=7,
                           dtype=np.float64)

loss, grads = model.loss(X, y)
for param_name in sorted(grads):
    f = lambda _: model.loss(X, y)[0]
    param_grad_num = eval_numerical_gradient(f, model.params[param_name], verbose=False,
    , h=1e-6)
    e = rel_error(param_grad_num, grads[param_name])
    print('{} max relative error: {}'.format(param_name, rel_error(param_grad_num, grads[param_name])))
```

```
W1 max relative error: 0.0016302996569852294
W2 max relative error: 0.0023566703879334696
W3 max relative error: 6.986613824440613e-05
b1 max relative error: 4.145999107053428e-05
b2 max relative error: 4.81077192964982e-07
b3 max relative error: 5.006131294224136e-07
```

Overfit small dataset

To check your CNN implementation, let's overfit a small dataset.

In [5]:

```
num_train = 100
small_data = {
    'X_train': data['X_train'][:num_train],
    'y_train': data['y_train'][:num_train],
    'X_val': data['X_val'],
    'y_val': data['y_val'],
}

model = ThreeLayerConvNet(weight_scale=1e-2)


solver = Solver(model, small_data,
                 num_epochs=10, batch_size=50,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 verbose=True, print_every=1)
solver.train()
```

```
(Iteration 1 / 20) loss: 2.320366
(Epoch 0 / 10) train acc: 0.220000; val_acc: 0.131000
(Iteration 2 / 20) loss: 3.262984
(Epoch 1 / 10) train acc: 0.070000; val_acc: 0.087000
(Iteration 3 / 20) loss: 3.300851
(Iteration 4 / 20) loss: 3.196094
(Epoch 2 / 10) train acc: 0.340000; val_acc: 0.159000
(Iteration 5 / 20) loss: 2.708954
(Iteration 6 / 20) loss: 2.365591
(Epoch 3 / 10) train acc: 0.320000; val_acc: 0.155000
(Iteration 7 / 20) loss: 2.398544
(Iteration 8 / 20) loss: 1.720295
(Epoch 4 / 10) train acc: 0.500000; val_acc: 0.206000
(Iteration 9 / 20) loss: 1.701536
(Iteration 10 / 20) loss: 1.752474
(Epoch 5 / 10) train acc: 0.500000; val_acc: 0.180000
(Iteration 11 / 20) loss: 1.587625
(Iteration 12 / 20) loss: 1.573282
(Epoch 6 / 10) train acc: 0.680000; val_acc: 0.222000
(Iteration 13 / 20) loss: 1.208725
(Iteration 14 / 20) loss: 1.107889
(Epoch 7 / 10) train acc: 0.720000; val_acc: 0.191000
(Iteration 15 / 20) loss: 1.209453
(Iteration 16 / 20) loss: 1.142800
(Epoch 8 / 10) train acc: 0.810000; val_acc: 0.170000
(Iteration 17 / 20) loss: 0.803844
(Iteration 18 / 20) loss: 0.726849
(Epoch 9 / 10) train acc: 0.790000; val_acc: 0.209000
(Iteration 19 / 20) loss: 0.820536
(Iteration 20 / 20) loss: 0.684563
(Epoch 10 / 10) train acc: 0.850000; val_acc: 0.226000
```

In [6]:

```
plt.subplot(2, 1, 1)
plt.plot(solver.loss_history, 'o')
plt.xlabel('iteration')
plt.ylabel('loss')

plt.subplot(2, 1, 2)
plt.plot(solver.train_acc_history, '-o')
plt.plot(solver.val_acc_history, '-o')
plt.legend(['train', 'val'], loc='upper left')
plt.xlabel('epoch')
plt.ylabel('accuracy')
plt.show()
```



Train the network

Now we train the 3 layer CNN on CIFAR-10 and assess its accuracy.

In [7]:

```
model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001)

solver = Solver(model, data,
                 num_epochs=1, batch_size=50,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 verbose=True, print_every=20)
solver.train()
```

```
(Iteration 1 / 980) loss: 2.304500
(Epoch 0 / 1) train acc: 0.109000; val_acc: 0.094000
(Iteration 21 / 980) loss: 2.195101
(Iteration 41 / 980) loss: 2.122926
(Iteration 61 / 980) loss: 1.562842
(Iteration 81 / 980) loss: 1.983982
(Iteration 101 / 980) loss: 1.515928
(Iteration 121 / 980) loss: 1.903064
(Iteration 141 / 980) loss: 1.852185
(Iteration 161 / 980) loss: 1.811960
(Iteration 181 / 980) loss: 1.776521
(Iteration 201 / 980) loss: 1.688216
(Iteration 221 / 980) loss: 1.721511
(Iteration 241 / 980) loss: 1.508774
(Iteration 261 / 980) loss: 1.708594
(Iteration 281 / 980) loss: 1.578662
(Iteration 301 / 980) loss: 1.765830
(Iteration 321 / 980) loss: 1.810934
(Iteration 341 / 980) loss: 1.462444
(Iteration 361 / 980) loss: 1.745883
(Iteration 381 / 980) loss: 1.627576
(Iteration 401 / 980) loss: 1.450037
(Iteration 421 / 980) loss: 1.854767
(Iteration 441 / 980) loss: 1.438163
(Iteration 461 / 980) loss: 1.475581
(Iteration 481 / 980) loss: 1.598954
(Iteration 501 / 980) loss: 1.333409
(Iteration 521 / 980) loss: 1.885802
(Iteration 541 / 980) loss: 1.927320
(Iteration 561 / 980) loss: 1.614946
(Iteration 581 / 980) loss: 1.592355
(Iteration 601 / 980) loss: 1.563484
(Iteration 621 / 980) loss: 1.641046
(Iteration 641 / 980) loss: 1.491252
(Iteration 661 / 980) loss: 1.626638
(Iteration 681 / 980) loss: 1.766847
(Iteration 701 / 980) loss: 1.713099
(Iteration 721 / 980) loss: 1.379815
(Iteration 741 / 980) loss: 1.478084
(Iteration 761 / 980) loss: 1.551079
(Iteration 781 / 980) loss: 1.569401
(Iteration 801 / 980) loss: 1.492906
(Iteration 821 / 980) loss: 2.004961
(Iteration 841 / 980) loss: 1.757052
(Iteration 861 / 980) loss: 1.643516
(Iteration 881 / 980) loss: 1.281183
(Iteration 901 / 980) loss: 1.606838
(Iteration 921 / 980) loss: 1.470896
(Iteration 941 / 980) loss: 1.391651
(Iteration 961 / 980) loss: 1.371272
(Epoch 1 / 1) train acc: 0.447000; val_acc: 0.449000
```

Get > 65% validation accuracy on CIFAR-10.

In the last part of the assignment, we'll now ask you to train a CNN to get better than 65% validation accuracy on CIFAR-10.

Things you should try:

- Filter size: Above we used 7x7; but VGGNet and onwards showed stacks of 3x3 filters are good.
- Number of filters: Above we used 32 filters. Do more or fewer do better?
- Batch normalization: Try adding spatial batch normalization after convolution layers and vanilla batch normalization after affine layers. Do your networks train faster?
- Network architecture: Can a deeper CNN do better? Consider these architectures:
 - [conv-relu-pool]xN - conv - relu - [affine]xM - [softmax or SVM]
 - [conv-relu-pool]xN - [affine]xM - [softmax or SVM]
 - [conv-relu-conv-relu-pool]xN - [affine]xM - [softmax or SVM]

Tips for training

For each network architecture that you try, you should tune the learning rate and regularization strength. When doing this there are a couple important things to keep in mind:

- If the parameters are working well, you should see improvement within a few hundred iterations
- Remember the coarse-to-fine approach for hyperparameter tuning: start by testing a large range of hyperparameters for just a few training iterations to find the combinations of parameters that are working at all.
- Once you have found some sets of parameters that seem to work, search more finely around these parameters. You may need to train for more epochs.

In [4]:

```
# ===== #
# YOUR CODE HERE:
#   Implement a CNN to achieve greater than 65% validation accuracy
#   on CIFAR-10.
# ===== #

model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001, filter_size=3)

solver = Solver(model, data,
                num_epochs=7, batch_size=256,
                update_rule='adam',
                optim_config={
                    'learning_rate': 1e-3,
                }, lr_decay=0.95,
                verbose=True, print_every=20)

solver.train()

# ===== #
# END YOUR CODE HERE
# ===== #
```

```
(Iteration 1 / 1337) loss: 2.304623
(Epoch 0 / 7) train acc: 0.104000; val_acc: 0.120000
(Iteration 21 / 1337) loss: 1.902956
(Iteration 41 / 1337) loss: 1.780110
(Iteration 61 / 1337) loss: 1.541009
(Iteration 81 / 1337) loss: 1.429866
(Iteration 101 / 1337) loss: 1.544923
(Iteration 121 / 1337) loss: 1.392739
(Iteration 141 / 1337) loss: 1.383102
(Iteration 161 / 1337) loss: 1.359704
(Iteration 181 / 1337) loss: 1.251945
(Epoch 1 / 7) train acc: 0.584000; val_acc: 0.545000
(Iteration 201 / 1337) loss: 1.355810
(Iteration 221 / 1337) loss: 1.354795
(Iteration 241 / 1337) loss: 1.297572
(Iteration 261 / 1337) loss: 1.185972
(Iteration 281 / 1337) loss: 1.149810
(Iteration 301 / 1337) loss: 1.253490
(Iteration 321 / 1337) loss: 1.112667
(Iteration 341 / 1337) loss: 1.304045
(Iteration 361 / 1337) loss: 1.179771
(Iteration 381 / 1337) loss: 1.095352
(Epoch 2 / 7) train acc: 0.632000; val_acc: 0.581000
(Iteration 401 / 1337) loss: 1.269374
(Iteration 421 / 1337) loss: 1.182422
(Iteration 441 / 1337) loss: 1.227449
(Iteration 461 / 1337) loss: 1.076570
(Iteration 481 / 1337) loss: 1.091211
(Iteration 501 / 1337) loss: 1.062462
(Iteration 521 / 1337) loss: 1.092874
(Iteration 541 / 1337) loss: 1.040742
(Iteration 561 / 1337) loss: 1.105746
(Epoch 3 / 7) train acc: 0.665000; val_acc: 0.610000
(Iteration 581 / 1337) loss: 1.085435
(Iteration 601 / 1337) loss: 1.026023
(Iteration 621 / 1337) loss: 1.029757
(Iteration 641 / 1337) loss: 0.852888
(Iteration 661 / 1337) loss: 0.987487
(Iteration 681 / 1337) loss: 1.097535
(Iteration 701 / 1337) loss: 0.954538
(Iteration 721 / 1337) loss: 0.956822
(Iteration 741 / 1337) loss: 0.979291
(Iteration 761 / 1337) loss: 0.877921
(Epoch 4 / 7) train acc: 0.698000; val_acc: 0.625000
(Iteration 781 / 1337) loss: 1.037775
(Iteration 801 / 1337) loss: 0.889407
(Iteration 821 / 1337) loss: 0.921538
(Iteration 841 / 1337) loss: 0.939175
(Iteration 861 / 1337) loss: 0.772293
(Iteration 881 / 1337) loss: 0.797997
(Iteration 901 / 1337) loss: 0.950400
(Iteration 921 / 1337) loss: 0.945413
(Iteration 941 / 1337) loss: 0.761205
(Epoch 5 / 7) train acc: 0.765000; val_acc: 0.637000
(Iteration 961 / 1337) loss: 0.827479
(Iteration 981 / 1337) loss: 0.876402
(Iteration 1001 / 1337) loss: 0.801727
(Iteration 1021 / 1337) loss: 0.928967
(Iteration 1041 / 1337) loss: 0.721841
(Iteration 1061 / 1337) loss: 0.793211
(Iteration 1081 / 1337) loss: 0.725237
```

```
(Iteration 1101 / 1337) loss: 0.783221
(Iteration 1121 / 1337) loss: 0.746461
(Iteration 1141 / 1337) loss: 0.828410
(Epoch 6 / 7) train acc: 0.802000; val_acc: 0.648000
(Iteration 1161 / 1337) loss: 0.737691
(Iteration 1181 / 1337) loss: 0.731964
(Iteration 1201 / 1337) loss: 0.690352
(Iteration 1221 / 1337) loss: 0.839792
(Iteration 1241 / 1337) loss: 0.713802
(Iteration 1261 / 1337) loss: 0.685698
(Iteration 1281 / 1337) loss: 0.663506
(Iteration 1301 / 1337) loss: 0.716763
(Iteration 1321 / 1337) loss: 0.834270
(Epoch 7 / 7) train acc: 0.808000; val_acc: 0.661000
```



```

1  import numpy as np
2  from nn1.layers import *
3  import pdb
4
5  """
6  This code was originally written for CS 231n at Stanford University
7  (cs231n.stanford.edu). It has been modified in various areas for use in the
8  ECE 239AS class at UCLA. This includes the descriptions of what code to
9  implement as well as some slight potential changes in variable names to be
10 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14
15 def conv_forward_naive(x, w, b, conv_param):
16     """
17     A naive implementation of the forward pass for a convolutional layer.
18
19     The input consists of N data points, each with C channels, height H and width
20     W. We convolve each input with F different filters, where each filter spans
21     all C channels and has height HH and width HH.
22
23     Input:
24     - x: Input data of shape (N, C, H, W)
25     - w: Filter weights of shape (F, C, HH, WW)
26     - b: Biases, of shape (F,)
27     - conv_param: A dictionary with the following keys:
28         - 'stride': The number of pixels between adjacent receptive fields in the
29           horizontal and vertical directions.
30         - 'pad': The number of pixels that will be used to zero-pad the input.
31
32     Returns a tuple of:
33     - out: Output data, of shape (N, F, H', W') where H' and W' are given by
34         H' = 1 + (H + 2 * pad - HH) / stride
35         W' = 1 + (W + 2 * pad - WW) / stride
36     - cache: (x, w, b, conv_param)
37     """
38     out = None
39     pad = conv_param['pad']
40     stride = conv_param['stride']
41
42     # ===== #
43     # YOUR CODE HERE:
44     # Implement the forward pass of a convolutional neural network.
45     # Store the output as 'out'.
46     # Hint: to pad the array, you can use the function np.pad.
47     # ===== #
48
49     N, C, H, W = x.shape
50     F, C, HH, WW = w.shape
51
52     # Add padding to each image
53     x_pad = np.pad(x, ((0,), (0,), (pad,), (pad,)), 'constant')
54     # Size of the output
55     Hh = 1 + int((H + 2 * pad - HH) / stride)
56     Hw = 1 + int((W + 2 * pad - WW) / stride)
57
58     out = np.zeros((N, F, Hh, Hw))
59
60     for n in range(N): # First, iterate over all the images
61         for f in range(F): # Second, iterate over all the kernels
62             for k in range(Hh):
63                 for l in range(Hw):
64                     out[n, f, k, l] = np.sum(
65                         x_pad[n, :, k * stride:k * stride + HH, l * stride:l * stride + WW] *
66                         w[f, :]) + b[f]

```

```

67
68
69
70 # ===== #
71 # END YOUR CODE HERE
72 # ===== #
73
74 cache = (x, w, b, conv_param)
75 return out, cache
76
77
78 def conv_backward_naive(dout, cache):
79     """
80     A naive implementation of the backward pass for a convolutional layer.
81
82     Inputs:
83     - dout: Upstream derivatives.
84     - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive
85
86     Returns a tuple of:
87     - dx: Gradient with respect to x
88     - dw: Gradient with respect to w
89     - db: Gradient with respect to b
90     """
91     dx, dw, db = None, None, None
92
93     N, F, out_height, out_width = dout.shape
94     x, w, b, conv_param = cache
95
96     stride, pad = [conv_param['stride'], conv_param['pad']]
97     xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
98     num_filts, _, f_height, f_width = w.shape
99
100 # ===== #
101 # YOUR CODE HERE:
102 # Implement the backward pass of a convolutional neural network.
103 # Calculate the gradients: dx, dw, and db.
104 # ===== #
105
106 N, F, H1, W1 = dout.shape
107 x, w, b, conv_param = cache
108 N, C, H, W = x.shape
109 HH = w.shape[2]
110 WW = w.shape[3]
111 stride = conv_param['stride']
112 pad = conv_param['pad']
113
114
115 dx, dw, db = np.zeros_like(x), np.zeros_like(w), np.zeros_like(b)
116 x_pad = np.pad(x, [(0,0), (0,0), (pad,pad), (pad,pad)], 'constant')
117 dx_pad = np.pad(dx, [(0,0), (0,0), (pad,pad), (pad,pad)], 'constant')
118 db = np.sum(np.sum(np.sum(dout, axis=0),axis=1),axis=1)
119
120 for n in np.arange(N):
121     for f in np.arange(F):
122         for i in np.arange(H1):
123             for j in np.arange(W1):
124                 # Window we want to apply the respective f th filter over (C, HH, WW)
125                 x_window = x_pad[n, :, i * stride : i * stride + HH, j * stride : j * stride
126                     + WW]
127
128                 dw[f] += x_window * dout[n, f, i, j]
129
130                 dx_pad[n, :, i * stride : i * stride + HH, j * stride : j * stride + WW] +=
131                     w[f] * dout[n, f, i, j]
132
133 dx = dx_pad[:, :, pad:pad+H, pad:pad+W]

```

```

132
133 # ===== #
134 # END YOUR CODE HERE
135 # ===== #
136
137 return dx, dw, db
138
139
140 def max_pool_forward_naive(x, pool_param):
141     """
142     A naive implementation of the forward pass for a max pooling layer.
143
144     Inputs:
145     - x: Input data, of shape (N, C, H, W)
146     - pool_param: dictionary with the following keys:
147       - 'pool_height': The height of each pooling region
148       - 'pool_width': The width of each pooling region
149       - 'stride': The distance between adjacent pooling regions
150
151     Returns a tuple of:
152     - out: Output data
153     - cache: (x, pool_param)
154     """
155     out = None
156
157     # ===== #
158     # YOUR CODE HERE:
159     # Implement the max pooling forward pass.
160     # ===== #
161
162     Hp = pool_param['pool_height']
163     Wp = pool_param['pool_width']
164     S = pool_param['stride']
165     N, C, H, W = x.shape
166     H1 = int((H - Hp) / S) + 1
167     W1 = int((W - Wp) / S) + 1
168
169     out = np.zeros((N, C, H1, W1))
170     for n in range(N):
171         for c in range(C):
172             for k in range(H1):
173                 for l in range(W1):
174                     out[n, c, k, l] = np.max(
175                         x[n, c, k * S:k * S + Hp, l * S:l * S + Wp])
176
177     # ===== #
178     # END YOUR CODE HERE
179     # ===== #
180
181     cache = (x, pool_param)
182     return out, cache
183
184 def max_pool_backward_naive(dout, cache):
185     """
186     A naive implementation of the backward pass for a max pooling layer.
187
188     Inputs:
189     - dout: Upstream derivatives
190     - cache: A tuple of (x, pool_param) as in the forward pass.
191
192     Returns:
193     - dx: Gradient with respect to x
194     """
195     dx = None
196     x, pool_param = cache
197     pool_height, pool_width, stride = pool_param['pool_height'],
    pool_param['pool_width'], pool_param['stride']

```

```

198 # ===== #
199 # YOUR CODE HERE:
200 #   Implement the max pooling backward pass.
201 # ===== #
202
203 N, C, H, W = x.shape
204 H1 = int((H - pool_height) / stride) + 1
205 W1 = int((W - pool_width) / stride) + 1
206
207 dx = np.zeros((N, C, H, W))
208 for nprime in range(N):
209     for cprime in range(C):
210         for k in range(H1):
211             for l in range(W1):
212                 x_pooling = x[nprime, cprime, k *
213                               stride:k * stride + pool_height, l *
214                               stride + pool_width]
215                 maxi = np.max(x_pooling)
216                 x_mask = x_pooling == maxi
217                 dx[nprime, cprime, k * stride:k * stride + pool_height, l *
218                   stride + pool_width] += dout[nprime, cprime, k, l] * x_mask
219
220 return dx
221
222 # ===== #
223 # END YOUR CODE HERE
224 # ===== #
225
226 return dx
227
228 def spatial_batchnorm_forward(x, gamma, beta, bn_param):
229     """
230     Computes the forward pass for spatial batch normalization.
231
232     Inputs:
233     - x: Input data of shape (N, C, H, W)
234     - gamma: Scale parameter, of shape (C,)
235     - beta: Shift parameter, of shape (C,)
236     - bn_param: Dictionary with the following keys:
237         - mode: 'train' or 'test'; required
238         - eps: Constant for numeric stability
239         - momentum: Constant for running mean / variance. momentum=0 means that
240           old information is discarded completely at every time step, while
241           momentum=1 means that new information is never incorporated. The
242           default of momentum=0.9 should work well in most situations.
243         - running_mean: Array of shape (D,) giving running mean of features
244         - running_var: Array of shape (D,) giving running variance of features
245
246     Returns a tuple of:
247     - out: Output data, of shape (N, C, H, W)
248     - cache: Values needed for the backward pass
249     """
250     out, cache = None, None
251
252     # ===== #
253     # YOUR CODE HERE:
254     #   Implement the spatial batchnorm forward pass.
255     #
256     #   You may find it useful to use the batchnorm forward pass you
257     #   implemented in HW #4.
258     # ===== #
259
260 N, C, H, W = x.shape
261 out, cache = batchnorm_forward(x.swapaxes(0, 1).reshape(C, N * H * W).T, gamma, beta,
                                bn_param)
262 out = out.T.reshape(C, N, H, W).swapaxes(0, 1)

```

```

262 # ===== #
263 # END YOUR CODE HERE
264 # ===== #
265
266 return out, cache
267
268
269 def spatial_batchnorm_backward(dout, cache):
270     """
271     Computes the backward pass for spatial batch normalization.
272
273     Inputs:
274     - dout: Upstream derivatives, of shape (N, C, H, W)
275     - cache: Values from the forward pass
276
277     Returns a tuple of:
278     - dx: Gradient with respect to inputs, of shape (N, C, H, W)
279     - dgamma: Gradient with respect to scale parameter, of shape (C,)
280     - dbeta: Gradient with respect to shift parameter, of shape (C,)
281     """
282     dx, dgamma, dbeta = None, None, None
283
284     # ===== #
285     # YOUR CODE HERE:
286     # Implement the spatial batchnorm backward pass.
287     #
288     # You may find it useful to use the batchnorm forward pass you
289     # implemented in HW #4.
290     # ===== #
291
292     N,C,H,W = dout.shape
293     dx, dgamma, dbeta = batchnorm_backward(dout.swapaxes(0,1).reshape(C,-1).T, cache)
294     dx = dx.T.reshape(C,N,H,W).swapaxes(0,1)
295
296     # ===== #
297     # END YOUR CODE HERE
298     # ===== #
299
300     return dx, dgamma, dbeta

```

```

1  import numpy as np
2
3  from nn1.layers import *
4  from nn1.conv_layers import *
5  from cs231n.fast_layers import *
6  from nn1.layer_utils import *
7  from nn1.conv_layer_utils import *
8
9  import pdb
10
11  """
12  This code was originally written for CS 231n at Stanford University
13  (cs231n.stanford.edu). It has been modified in various areas for use in the
14  ECE 239AS class at UCLA. This includes the descriptions of what code to
15  implement as well as some slight potential changes in variable names to be
16  consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
17  permission to use this code. To see the original version, please visit
18  cs231n.stanford.edu.
19  """
20
21  class ThreeLayerConvNet(object):
22      """
23      A three-layer convolutional network with the following architecture:
24
25      conv - relu - 2x2 max pool - affine - relu - affine - softmax
26
27      The network operates on minibatches of data that have shape (N, C, H, W)
28      consisting of N images, each with height H and width W and with C input
29      channels.
30      """
31
32      def __init__(self, input_dim=(3, 32, 32), num_filters=32, filter_size=7,
33                  hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=0.0,
34                  dtype=np.float32, use_batchnorm=False):
35          """
36          Initialize a new network.
37
38          Inputs:
39          - input_dim: Tuple (C, H, W) giving size of input data
40          - num_filters: Number of filters to use in the convolutional layer
41          - filter_size: Size of filters to use in the convolutional layer
42          - hidden_dim: Number of units to use in the fully-connected hidden layer
43          - num_classes: Number of scores to produce from the final affine layer.
44          - weight_scale: Scalar giving standard deviation for random initialization
45            of weights.
46          - reg: Scalar giving L2 regularization strength
47          - dtype: numpy datatype to use for computation.
48          """
49          self.use_batchnorm = use_batchnorm
50          self.params = {}
51          self.reg = reg
52          self.dtype = dtype
53
54
55          # ===== #
56          # YOUR CODE HERE:
57          #   Initialize the weights and biases of a three layer CNN. To initialize:
58          #   - the biases should be initialized to zeros.
59          #   - the weights should be initialized to a matrix with entries
60          #     drawn from a Gaussian distribution with zero mean and
61          #     standard deviation given by weight_scale.
62          # ===== #
63
64          C, H, W = input_dim
65
66          height = int((H-2)/2+1)
67          width = int((W-2)/2+1)

```

```

68
69 weight_dimensions = [(num_filters, C, filter_size,
70 filter_size), (height*width*num_filters, hidden_dim), (hidden_dim, num_classes)]
71
72 bias_dimensions = [num_filters, hidden_dim, num_classes]
73
74
75 for i in np.arange(1,4):
76     self.params['W%d' %i] = np.random.normal(loc=0.0,
77     scale=weight_scale,size=weight_dimensions[i-1]) #weights are normall distributed
78     self.params['b%d' %i] = np.zeros(bias_dimensions[i-1])
79
80     # ===== #
81     # END YOUR CODE HERE
82     # ===== #
83
84     for k, v in self.params.items():
85         self.params[k] = v.astype(dtype)
86
87
88 def loss(self, X, y=None):
89     """
90     Evaluate loss and gradient for the three-layer convolutional network.
91
92     Input / output: Same API as TwoLayerNet in fc_net.py.
93     """
94     W1, b1 = self.params['W1'], self.params['b1']
95     W2, b2 = self.params['W2'], self.params['b2']
96     W3, b3 = self.params['W3'], self.params['b3']
97
98     # pass conv_param to the forward pass for the convolutional layer
99
100     filter_size = W1.shape[2]
101
102     conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}
103
104     # pass pool_param to the forward pass for the max-pooling layer
105     pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}
106
107     scores = None
108
109     # ===== #
110     # YOUR CODE HERE:
111     # Implement the forward pass of the three layer CNN. Store the output
112     # scores as the variable "scores".
113     # ===== #
114
115     conv_out, conv_cache = conv_forward_fast(X, W1, b1, conv_param)
116     conv_relu, conv_relu_cache = relu_forward(conv_out)
117     pool_out, pool_cache = max_pool_forward_fast(conv_relu, pool_param)
118     affine_out, affine_cache = affine_forward(pool_out, W2, b2)
119     affine_relu, affine_relu_cache = relu_forward(affine_out)
120     scores, scores_cache = affine_forward(affine_relu, W3, b3)
121
122     # ===== #
123     # END YOUR CODE HERE
124     # ===== #
125
126     if y is None:
127         return scores
128
129     loss, grads = 0, {}
130     # ===== #
131     # YOUR CODE HERE:
132     # Implement the backward pass of the three layer CNN. Store the grads

```

```

133     # in the grads dictionary, exactly as before (i.e., the gradient of
134     # self.params[k] will be grads[k]). Store the loss as "loss", and
135     # don't forget to add regularization on ALL weight matrices.
136     # ===== #
137
138     loss, output_derivative = softmax_loss(scores, y)
139     loss += 0.5*self.reg*(np.sum(W1*W1)+np.sum(W2*W2)+np.sum(W3*W3))
140     affine_derivative, grads['W3'], grads['b3'] = affine_backward(output_derivative,
141     scores_cache)
142     relu_derivative = relu_backward(affine_derivative, affine_relu_cache)
143     affine_derivative, grads['W2'], grads['b2'] = affine_backward(relu_derivative,
144     affine_cache)
145
146     pool_derivative = max_pool_backward_fast(affine_derivative, pool_cache)
147     x_derivative = relu_backward(pool_derivative, conv_relu_cache)
148     x_derivative, grads['W1'], grads['b1'] = conv_backward_fast(x_derivative, conv_cache)
149
150     grads['W3'] += self.reg*W3
151     grads['W2'] += self.reg*W2
152     grads['W1'] += self.reg*W1
153
154
155     # ===== #
156     # END YOUR CODE HERE
157     # ===== #
158
159     return loss, grads
160
161
162 pass
163

```